Detecting Heavy Drinking with Smartphone Accelerometers

Grace Tang
Springboard DSC - Capstone 2

The Problem with Alcohol

More than 95,000 people die from excessive alcohol use in the U.S. each year

cdc.gov/alcohol

The Solution: Smartphones and Algorithms

- Can we predict incidences of heavy drinking using smartphone accelerometer data?
- Many applications:
  - Just-In-Time Adaptive Interventions (JITAIs)
  - Law Enforcement
 - Medical Intervention

The Data: Accelerometers and TAC

- **13 anonymous persons**
  - 11 iPhones
  - 2 Android phones
- **Accelerometer data**
  - Sampling rate: 40 Hz
  - 14 million rows
- **Transdermal Alcohol Content (TAC)**
  - Sampling rate: per 30 min
  - 600 rows

	time	pid	x	y	z
0	0	JB3156	0.0000	0.0000	0.0000
1	0	CC6740	0.0000	0.0000	0.0000
2	1493733882409	SA0297	0.0758	0.0273	-0.0102
3	1493733882455	SA0297	-0.0359	0.0794	0.0037
4	1493733882500	SA0297	-0.2427	-0.0861	-0.0163

|	TAC Level	Time
0	0.0	2017-05-02 10:36:54
1	0.0	2017-05-02 11:09:57
2	0.0	2017-05-02 11:15:27
3	0.0	2017-05-02 11:20:57
4	0.0	2017-05-02 11:26:26

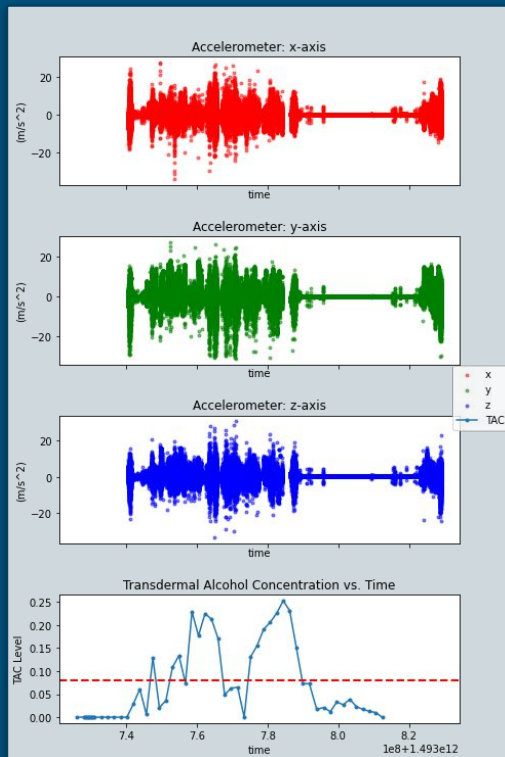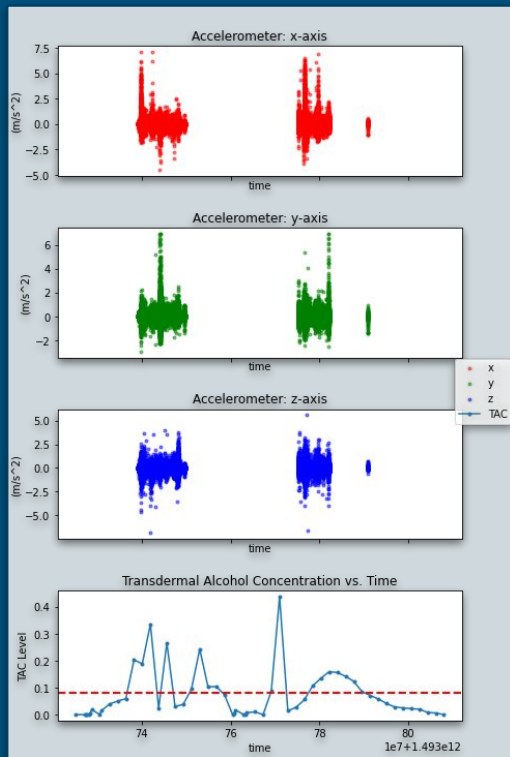# Bifurcating the Dataset: 2 Ways to Impute TAC

- Smooth with TMA
- Then scipy interpolate
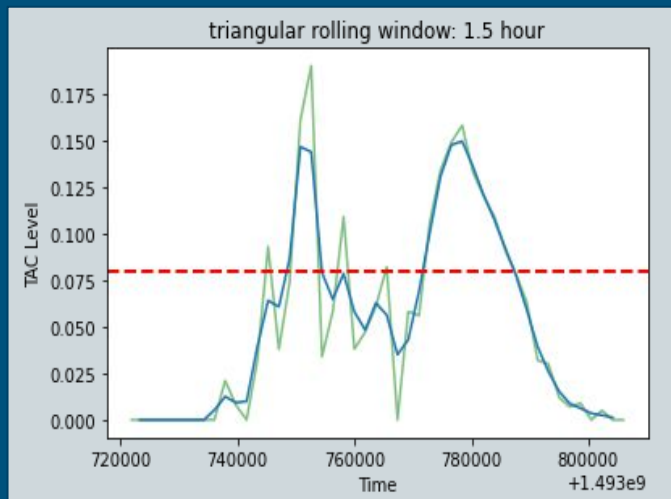
- Preserve the raw datapoints
- Then scipy interpolate

Accelerometers: Noisy Data and Missing Data

Featurization: Making Meaning out of Noise

- Before featurization: 3 features (x, y, z) & 1 target (TAC)
- Pandas rolling() method, 10-second windows
- After featurization: 306 new features
  - Mean, Standard Deviation, Variance, Median, Max & Min (of raw & absolute signal), Skew, Kurtosis, Zero Crossing Rate, Gait stretch, Number of steps, Step Time, Root Mean Squared, Average Resultant Acceleration, partial derivatives of each axes
  - Jerk (the derivative of acceleration) + above
  - Snap (the derivative of jerk) + above
  - The standard deviation of all of the above features

No Immediate Patterns

Boxplot of X-axis Features for a Single Participant

Sequential Processing is slow…

- 23 million rows to featurize
  - 13 dataframes (1 per person) x 2 (raw vs smoothed)
- Time taken to featurize:
  - 10,000 rows - 127 seconds
  - 447423 rows - 3414 seconds
  - (23 million rows - 48.2 hours!)

Using Multiprocessing to Parallelize

- Python's multiprocessing package
- Challenges with Windows OS

- Result: 48732 seconds (13.5 hours)
  - 3.5x faster
  - Reduce from 2-day to overnight!

```python
import defs
if __name__ == '__main__':
    num_processors = 10
    p=Pool(processes = num_processors)
    output = p.map(defs.featurize_df,[i for i in df_list])
    for i in range(len(df_list)):
        output[i].to_csv('pid'+str(i+1)+'_featurized.csv')
        print('pid'+str(i+1)+'_featurized.csv')
```

Preprocessing Data

Reduce from 53.0 GB to 42.4 MB

- Keep only every 100th row
- Random select n = 400/combo
- Drop columns with > 100 NaNs
- Drop rows with NaNs

- Raw/interpolate dataframe
  - 9107 rows, 293 columns
- TMA-smoothed dataframe
  - 9486 rows, 287 columns
- 70/30 train/test split

Number of TAC events per Participant (tma-smoothed data)

Lots and Lots of Decision Trees!

Binary Classification Models:

- ○ Random Forest
- ○ AdaBoost
- ○ Gradient Boost
- ○ XGBoost

Hyperparameter Tuning: Random Search Cross Validation

- ○ XGBoost Best Params: {'n_estimators': 1000, 'max_depth': 15, 'learning_rate': 0.1}

Table of Results

Model	Data	Accuracy	F1 -Score	Precision	Recall	P-R AUC	ROC AUC	Runtime
RF (paper)	-	0.775	-	0.666	0.698	-	-	-
RF	raw	0.753	0.748	0.718	0.873	0.82	0.69	2.526
RF	tma	0.793	0.791	0.742	0.903	0.88	0.72	**2.692**
AdaBoost	raw	0.680	0.679	0.680	0.741	0.75	0.64	34.953
AdaBoost	tma	0.711	0.710	0.694	0.759	0.79	0.65	35.277
GradientBoost	raw	0.739	0.737	0.719	0.830	0.81	0.69	85.575
GradientBoost	tma	0.762	0.760	0.723	0.853	0.84	0.69	83.256
XGBoost	raw	0.864	0.863	0.836	0.922	0.94	0.81	4.655
XGBoost	tma	**0.894**	**0.893**	**0.875**	**0.920**	**0.96**	0.84	3.641
RF - tuned	tma	0.768	0.765	0.722	0.874	0.85	0.69	**0.528**
XGBoost - tuned	tma	**0.909**	**0.909**	**0.891**	**0.931**	**0.97**	0.86	32.856

Model Comparison: ROC and Precision-Recall
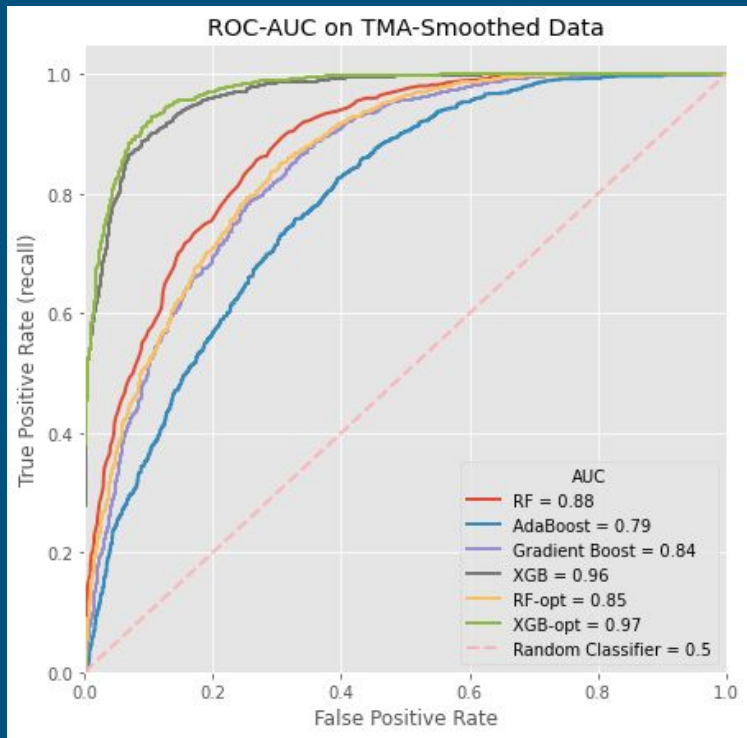
Feature Importance:
Standard Deviations of Other Features

Feature Importance: XGB Optimized

Takeaways: XGB is Best & Features are Good

- XGBoost: best model, but heavy.
    - Out-of-the-box XGB: 9x faster than tuned XGB
    - RF: much faster, but needs additional tuning
- Features:
    - Adding Jerk, Snap, and Standard Deviation of other features improves model prediction

Future Directions: Make it Lighter

- RF Classifier - additional tuning
- XGB Classifier - additional tuning, cap n_estimators $\leq$ 100
- Rewrite featurization step - feature reduction
- Alternative models:
  - ARIMA
  - Facebook Prophet

Thank you

Jackson A Killian (jkillian '@' g.harvard.edu, Harvard University); Danielle R Madden (University of Southern California); John Clapp (University of Southern California) for uploading this valuable dataset to the UCI Machine Learning Repository.

Chris Esposo for his guidance and being and awesome Springboard mentor!

Citations

Killian, J.A., Passino, K.M., Nandi, A., Madden, D.R. and Clapp, J., Learning to Detect Heavy Drinking Episodes Using Smartphone Accelerometer Data. In Proceedings of the 4th International Workshop on Knowledge Discovery in Healthcare Data co-located with the 28th International Joint Conference on Artificial Intelligence (IJCAI 2019) (pp. 35-42). [http://ceur-ws.org/Vol-2429/paper6.pdf]

Dataset:
http://archive.ics.uci.edu/ml/datasets/Bar+Crawl%3A+Detecting+Heavy+Drinking

Questions?